

r4nhldtj5

March 3, 2025

```
[5]: # Import all the Dependencies
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
[17]: # Data Collection and analysis
from google.colab import files
uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving insurance.csv to insurance (1).csv

```
[7]: # Load the csv dataset into a pandas dataframe
insurance_dataset = pd.read_csv("insurance.csv")
```

```
[8]: # Display the first five rows of a dataset
insurance_dataset.head()
```

```
[8]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
[9]: # Find the number of rows and columns
insurance_dataset.shape
```

```
[9]: (1338, 7)
```

```
[10]: # Getting some information about the dataset
insurance_dataset.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1338 non-null   int64
 1   sex         1338 non-null   object
 2   bmi         1338 non-null   float64
 3   children    1338 non-null   int64
 4   smoker      1338 non-null   object
 5   region      1338 non-null   object
 6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB

```

```

[11]: # We three categorical features
      # 1 whether a person is a smoker or non-smoker (smoker)
      # 2 whether the person is male or female (sex)
      # 3 The region
      # Charges is the target variable

```

```

[12]: # Checking for missing values
      insurance_dataset.isnull().sum()

```

```

[12]: age         0
      sex         0
      bmi         0
      children    0
      smoker      0
      region      0
      charges     0
      dtype: int64

```

```

[13]: # Data Analysis
      # Statistical measures about the dataset
      insurance_dataset.describe()

```

```

[13]:
count    1338.000000    1338.000000    1338.000000    1338.000000
mean      39.207025     30.663397      1.094918    13270.422265
std       14.049960      6.098187      1.205493    12110.011237
min       18.000000     15.960000      0.000000     1121.873900
25%       27.000000     26.296250      0.000000     4740.287150
50%       39.000000     30.400000      1.000000     9382.033000
75%       51.000000     34.693750      2.000000    16639.912515
max       64.000000     53.130000      5.000000    63770.428010

```

```
[14]: # Find the distribution of age value
sns.set()
plt.figure(figsize=(6,6))
sns.distplot(insurance_dataset['age'])
plt.title('Age Distribution')
plt.show()
```


<ipython-input-14-2e9101a5988b>:4: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(insurance_dataset['age'])
```



```
[20]: # Finding the distribution for gender column
plt.figure(figsize=(6,6))
sns.countplot(x='sex', data=insurance_dataset)
plt.title("Sex Distribution")
plt.show()
```



```
[21]: insurance_dataset['sex'].value_counts() -# Checking values for the gender and how it is distributed
```

```
[21]: sex
      male      676
      female    662
      Name: count, dtype: int64
```

```
[23]: # Finding the distribution of the BMI
      plt.figure(figsize=(6,6))
      sns.distplot(insurance_dataset['bmi'])
      plt.title('BML Distribution')
      plt.show()
```


<ipython-input-23-5572a4a62209>:3: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(insurance_dataset['bmi'])
```



```
[24]: # Normal BMI range for a person is 18,5-24,9  
      # if bml< 18,5 - person is underweight  
      # if bmi is > 24,9 - person is overweight
```

```
# Based on the graph above, a lot of people are overweight this can influence the insurance cost a person get. It might increase the insurance cost.
```

```
[25]: # Use count plot for finding the distribution of the Children column
plt.figure(figsize=(6,6))
sns.countplot(x='children', data=insurance_dataset)
plt.title('Children Distribution')
plt.show()
```




```
[26]: # Find the number of the Children using value counts
insurance_dataset['children'].value_counts()
```

```
[26]: children
0      574
1      324
2      240
3      157
4       25
5       18
```

Name: count, dtype: int64

```
[28]: # Finding the distribution for the smoker column
plt.figure(figsize=(6,6))
sns.countplot(x='smoker', data=insurance_dataset)
plt.title('Smoker Distribution')
plt.show()
```




```
[30]: insurance_dataset['smoker'].value_counts() - # the numerical value of the
      ↪ distribution of the smoker column
```

```
[30]: smoker
no      1064
yes      274
Name: count, dtype: int64
```

```
[33]: # Finding the distribution of the regions column
plt.figure(figsize=(6,6))
sns.countplot(x='region', data=insurance_dataset)
plt.title('Region Distribution')
```



```
plt.show()
```



```
[35]: insurance_dataset['region'].value_counts() # the numerical value of the column ↵  
      ↪ region
```

```
[35]: region  
      southeast    364  
      southwest    325  
      northwest    325  
      northeast    324  
      Name: count, dtype: int64
```

```
[37]: # Finding the distribution for the charges using distribution plot  
      plt.figure(figsize=(6,6))  
      sns.distplot(insurance_dataset['charges'])  
      plt.title('Charges Distribution')  
      plt.show()
```


<ipython-input-37-3b6775a2176c>:3: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(insurance_dataset['charges'])
```



```
[40]: # Data Preprocessing

# Encoding the categorical features

# encoding the sex column
insurance_dataset.replace({'sex':{'male':0, 'female':1}}, inplace=True)

# Encoding the smoker column
insurance_dataset.replace({'smoker':{'yes':0, 'no':1}}, inplace=True)
```

```
# Encoding the region column
insurance_dataset.replace({'region':{'southwest':0, 'northwest':1, 'southeast':
↪2, 'northeast':3}}, inplace=True)
```

```
<ipython-input-40-05afef55c299>:6: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain the
old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
insurance_dataset.replace({'sex':{'male':0, 'female':1}}, inplace=True)
<ipython-input-40-05afef55c299>:9: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain the
old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
insurance_dataset.replace({'smoker':{'yes':0, 'no':1}}, inplace=True)
<ipython-input-40-05afef55c299>:12: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To retain the
old behavior, explicitly call `result.infer_objects(copy=False)`. To opt-in to
the future behavior, set `pd.set_option('future.no_silent_downcasting', True)`
insurance_dataset.replace({'region':{'southwest':0,
'northwest':1, 'southeast':2, 'northeast':3}}, inplace=True)
```

```
[41]: print(insurance_dataset)
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	0	0	16884.92400
1	18	0	33.770	1	1	2	1725.55230
2	28	0	33.000	3	1	2	4449.46200
3	33	0	22.705	0	1	1	21984.47061
4	32	0	28.880	0	1	1	3866.85520
...
1333	50	0	30.970	3	1	1	10600.54830
1334	18	1	31.920	0	1	3	2205.98080
1335	18	1	36.850	0	1	2	1629.83350
1336	21	1	25.800	0	1	0	2007.94500
1337	61	1	29.070	0	0	1	29141.36030

[1338 rows x 7 columns]

```
[44]: # We have to split the features and target
X = insurance_dataset.drop(columns='charges', axis=1)
Y = insurance_dataset['charges']
print(X)
print(Y)
```

	age	sex	bmi	children	smoker	region
0	19	1	27.900	0	0	0
1	18	0	33.770	1	1	2
2	28	0	33.000	3	1	2

3	33	0	22.705	0	1	1
4	32	0	28.880	0	1	1
...
1333	50	0	30.970	3	1	1
1334	18	1	31.920	0	1	3
1335	18	1	36.850	0	1	2
1336	21	1	25.800	0	1	0
1337	61	1	29.070	0	0	1

[1338 rows x 6 columns]

0	16884.92400
1	1725.55230
2	4449.46200
3	21984.47061
4	3866.85520

...	...
1333	10600.54830
1334	2205.98080
1335	1629.83350
1336	2007.94500
1337	29141.36030

Name: charges, Length: 1338, dtype: float64

```
[45]: # Split the data into training data and testing data
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
↳ random_state = 2)
```

```
[46]: # Print the shape of this dataset
print(X.shape, X_train.shape, X_test.shape)
```

(1338, 6)		age	sex	bmi	children	smoker	region
882	21	1	22.135	0	1	3	
505	37	0	30.875	3	1	1	
798	58	1	33.100	0	1	0	
792	22	1	23.180	0	1	3	
201	48	1	32.230	1	1	2	
...
466	60	1	28.700	1	1	0	
299	48	1	28.880	1	1	1	
493	61	0	43.400	0	1	0	
527	51	1	25.800	1	1	0	
1192	58	1	32.395	1	1	3	

[1070 rows x 6 columns] (268, 6)

```
[47]: # Model training
# Using linear regression model
```

```
# Loading the linear regression model  
regressor = LinearRegression()
```

```
[48]: # Fit the model into X & Y  
regressor.fit(X_train,Y_train)
```

```
[48]: LinearRegression()
```

```
[49]: # Model Evaluation  
# Use R squared  
# prediction on training data  
training_data_prediction = regressor.predict(X_train)
```

```
[51]: # Finding the r squared value  
r2_train = metrics.r2_score(Y_train, training_data_prediction)  
print('R squared value:',r2_train)
```

R squared value: 0.751618558520502

```
[52]: # R 2 squared is closer to 1 which means the model performed well.
```

```
[53]: # Finding the r squared value on test data  
  
test_data_prediction = regressor.predict(X_test)  
r2_test = metrics.r2_score(Y_test, test_data_prediction)  
print('R squared value:', r2_test)
```

R squared value: 0.7428748503048913

```
[ ]: # Similiar to the score obtained on training data- this means the model did  
↪ well.
```

```
[ ]:
```